

Amazon Reviews of Musical Equipment

A NLP Project

Final Capstone Project

Intro

As a consumer in this day and age, I do most of my shopping online. This trend of doing my shopping online created a certain level of skepticism when I buy products due to a few bad experiences I have had with my purchases.

Due to this fact, product reviews are extremely important to me because they determine if I should purchase a product or not. Such product review data can create a huge impact on the sales of a product from a companies point of view.

Data;

The data was retrieved from <http://jmcauley.ucsd.edu/data/amazon/> by accessing the website. The variables from the dataset that would be utilized, are the reviews and the sentiment in the review (positive versus negative).

Use of Specialization;

Techniques intended to be used include use of several plotting techniques to understand the data, use NLP and NLTK (the use of text processing methods such as tokenizations, stop word removal, stemming and vectorizing text via term frequencies (TF) as well as the inverse document frequencies (TF-IDF)). The use of topic modelling would be done with Latent Dirichlet Allocation (LDA) and Sentiment Analysis. For evaluation, Recall and F1 score would be used in addition to Receiver Operating Characteristic score and Confusion Matrix

Product/Business Impact;

This project is valuable because it provides actionable insight on customers reactions to a product. It can assist in better marketing of the product and directed improvement to increase customer satisfaction, which in turn may provide a more favorable view of the product and increase sales. In addition, the model created would assist in predicting whether future reviews are positive or negative towards the product.

Project goals include;

- Properly cleaning the data and using visualization to understand the dataset,
- Applying at least 2 models to determine process with highest accuracy.
- Creating a Sentiment analysis and a predictive model.

Being able to analyse that data to determine consumers' sentiments of the products, the flaws and even the most enjoyed feature can allow a company to improve on the product, and market it better based on the top feature mentioned.

The Dataset

The data contains the following columns:

asin – ID of the product

helpful – helpfulness rating of the review, e.g. 2/3

overall – rating of the product

reviewText – text of the review

reviewTime – time of the review (raw)

reviewerID – ID of the reviewer

reviewerName – name of the reviewer

summary – summary of the review

unixReviewTime – time of the review (unix time)

Lets import some packages

```
1 # Import Packages
2 import numpy as np
3 import pandas as pd
4 import plotly.offline as py|
5 py.init_notebook_mode(connected=True)
6 import plotly.graph_objs as go
7 import plotly.tools as tls
8 from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
9 from sklearn.decomposition import NMF, LatentDirichletAllocation
10 from matplotlib import pyplot as plt
11 %matplotlib inline
12 import scipy
13 import seaborn as sns
14 import re
15
16 import warnings
17 warnings.filterwarnings('ignore')
```

Exploring the data

```
1 df = pd.read_json('Musical_Instruments_5.json', lines=True)
2 df.dropna()
3 df.head()
```

	asin	helpful	overall	reviewText	reviewTime	reviewerID	reviewerName	summary	unixReviewTime
0	1384719342	[0, 0]	5	Not much to write about here, but it does exac...	02 28, 2014	A2IBPI20UZIR0U	cassandra tu "Yeah, well, that's just like, u...	good	1393545600
1	1384719342	[13, 14]	5	The product does exactly as it should and is q...	03 16, 2013	A14VAT5EAX3D9S	Jake	Jake	1363392000
2	1384719342	[1, 1]	5	The primary job of this device is to block the...	08 28, 2013	A195EZSQDW3E21	Rick Bennette "Rick Bennette"	It Does The Job Well	1377648000
3	1384719342	[0, 0]	5	Nice windscreen protects my MXL mic and preven...	02 14, 2014	A2C00NNG1ZQQG2	RustyBill "Sunday Rocker"	GOOD WINDSCREEN FOR THE MONEY	1392336000
4	1384719342	[0, 0]	5	This pop filter is great. It looks and perform...	02 21, 2014	A94QU4C90B1AX	SEAN MASLANKA	No more pops when I record my vocals.	1392940800


```
1 # Start exploring the dataset. Lets take a look at the columns
2 print(df.columns)
```

```
Index(['asin', 'helpful', 'overall', 'reviewText', 'reviewTime', 'reviewerID',
       'reviewerName', 'summary', 'unixReviewTime'],
      dtype='object')
```

```
1 print("There are {} observations and {} features in this dataset. \n".format(df.shape[0],df.shape[1]))
2 # Data types
3 df.dtypes
```

There are 10261 observations and 9 features in this dataset.

```
asin      object
helpful    object
overall    int64
reviewText object
reviewTime object
reviewerID object
reviewerName object
summary    object
unixReviewTime int64
dtype: object
```

```
: 1 # Checking for Null columns
2 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10261 entries, 0 to 10260
Data columns (total 9 columns):
asin                10261 non-null object
helpful             10261 non-null object
overall            10261 non-null int64
reviewText          10261 non-null object
reviewTime          10261 non-null object
reviewerID          10261 non-null object
reviewerName        10234 non-null object
summary             10261 non-null object
unixReviewTime      10261 non-null int64
dtypes: int64(2), object(7)
memory usage: 721.6+ KB
```

```
: 1 # Groupby by rating
2 reviewText_ = df.groupby("overall")
3
4 # Summary statistic of all sentiments
5 reviewText_.describe()
```

```
:

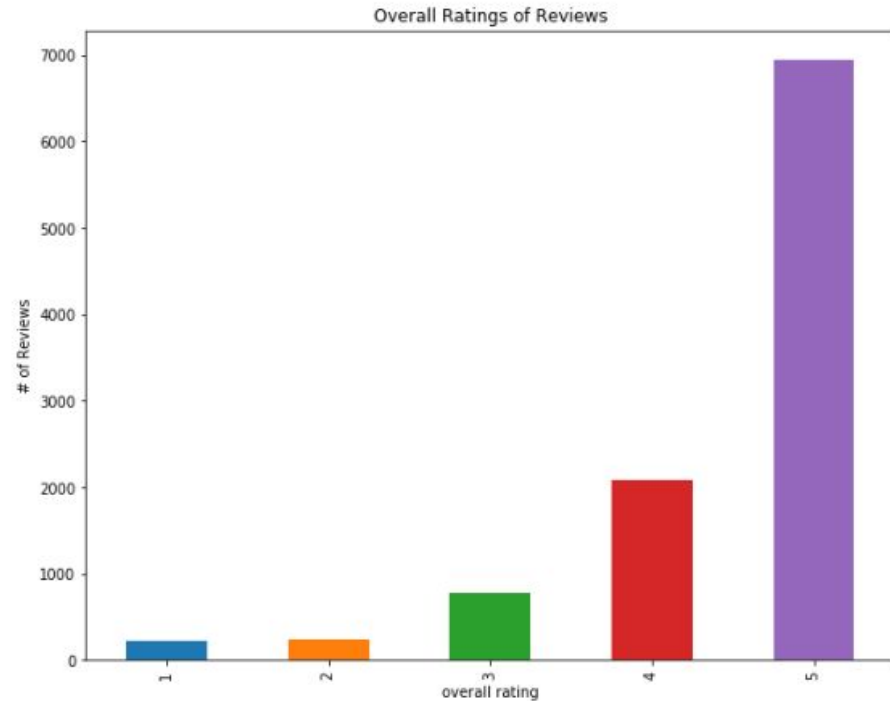
```

		count	mean	std	min	25%	50%	75%	unixReviewTime max
overall									
	1	217.0	1.363610e+09	3.693997e+07	1.141344e+09	1.347926e+09	1.370995e+09	1.390349e+09	1.405210e+09
	2	250.0	1.361242e+09	3.770940e+07	1.190678e+09	1.342116e+09	1.369872e+09	1.389506e+09	1.405210e+09
	3	772.0	1.361718e+09	3.633831e+07	1.161389e+09	1.343282e+09	1.369008e+09	1.389053e+09	1.405901e+09
	4	2084.0	1.359799e+09	3.914760e+07	1.095466e+09	1.342915e+09	1.369138e+09	1.388707e+09	1.405987e+09
	5	6938.0	1.360608e+09	3.757515e+07	1.096416e+09	1.343606e+09	1.367971e+09	1.388945e+09	1.405901e+09

Visualization of the review distribution based on the rating

(1 being the most negative and 5 being the most positive)

```
1 # Visualization of the review distribution based on the rating
2 df['overall'].value_counts().sort_values().plot(kind='bar', figsize=(10,8))
3 plt.title('Overall Ratings of Reviews')
4 plt.ylabel('# of Reviews')
5 plt.xlabel('overall rating')
6 plt.show()
```



```

1 # Drop missing values
2 df.dropna(inplace=True)
3
4 # Remove any 'neutral' ratings equal to 3
5 df = df[df['overall'] != 3]
6
7 # Encode 4s and 5s as 1 (rated positively)
8 # Encode 1s and 2s as 0 (rated negatively)
9 df['rating'] = np.where(df['overall'] > 3, 1, 0)
10 df.head()

```

	asin	helpful	overall	reviewText	reviewTime	reviewerID	reviewerName
0	1384719342	[0, 0]	5	Not much to write about here, but it does exac...	02 28, 2014	A2IBPI20UZIR0U	cassandra tu "Yea well, that's just lik...
1	1384719342	[13, 14]	5	The product does exactly as it should and is q...	03 16, 2013	A14VAT5EAX3D9S	Ja
2	1384719342	[1, 1]	5	The primary job of this device is to block the...	08 28, 2013	A195EZSQDW3E21	Rick Bennette "Ric Bennet
3	1384719342	[0, 0]	5	Nice windscreen protects my MXL mic and preven...	02 14, 2014	A2C00NNG1ZQQG2	RustyBill "Sund Rocke
4	1384719342	[0, 0]	5	This pop filter is great. It looks and perform...	02 21, 2014	A94QU4C90B1AX	SEAN MASLANH

Lets remove any neutral rating.

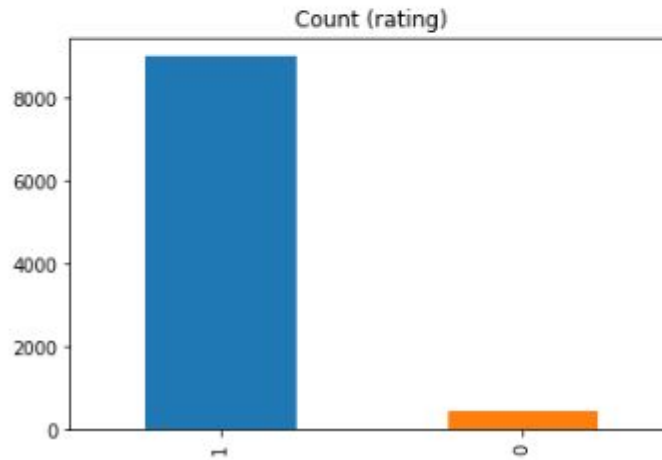
A rating of 3 would be considered a neutral rating.

- Ratings of 4 and 5 would be considered a positive rating
- Ratings of 1 and 2 would be considered a negative rating

Visualization of the positive versus negative reviews

```
1 # Visualization of the positive versus negative reviews
2
3 rating_count = df.rating.value_counts()
4 print('Negative Rating:', rating_count[0])
5 print('Positive Rating:', rating_count[1])
6 print('Proportion:', round(rating_count[0] / rating_count[1], 2), ': 1')
7
8 rating_count.plot(kind='bar', title='Count (rating)');
```

Negative Rating: 465
Positive Rating: 8998
Proportion: 0.05 : 1



Length of Reviews

Finding correlation between the length of the positive and negative reviews

```
1 def length(text):
2     '''a function which returns the length of text'''
3     return len(text)
4
5 # Apply the function to each review
6
7 df['length'] = df['reviewText'].apply(length)
8 df.head()
```

	asin	helpful	overall	reviewText	reviewTime	reviewerID	reviewerName	summary	unixReviewTime	rating	length
0	1384719342	[0, 0]	5	Not much to write about here, but it does exac...	02 28, 2014	A2IBPI20UZIR0U	cassandra tu "Yeah, well, that's just like, u...	good	1393545600	1	268
1	1384719342	[13, 14]	5	The product does exactly as it should and is q...	03 16, 2013	A14VAT5EAX3D9S	Jake	Jake	1363392000	1	544
2	1384719342	[1, 1]	5	The primary job of this device is to block the...	08 28, 2013	A195EZSQDW3E21	Rick Bennette "Rick Bennette"	It Does The Job Well	1377648000	1	436
3	1384719342	[0, 0]	5	Nice windscreen protects my MXL mic and preven...	02 14, 2014	A2C00NNG1ZQQG2	RustyBill "Sunday Rocker"	GOOD WINDSCREEN FOR THE MONEY	1392336000	1	206
4	1384719342	[0, 0]	5	This pop filter is great. It looks and perform...	02 21, 2014	A94QU4C90B1AX	SEAN MASLANKA	No more pops when I record my vocals.	1392940800	1	159

Average length of the Reviews

```
1 # Lets create a dataset for the positive and negative reviews
2
3 positive_review = df[df['rating'] == 1]
4 negative_review = df[df['rating'] == 0]
5
```

```
1 # Lets calculate the average length of the reviews
2
3 positive_review_mean = np.mean(positive_review['length'])
4 negative_review_mean = np.mean(negative_review['length'])
5 print('Average length of the Postive reviews:', positive_review_mean)
6 print('Average length of the Negative reviews:', negative_review_mean)
```

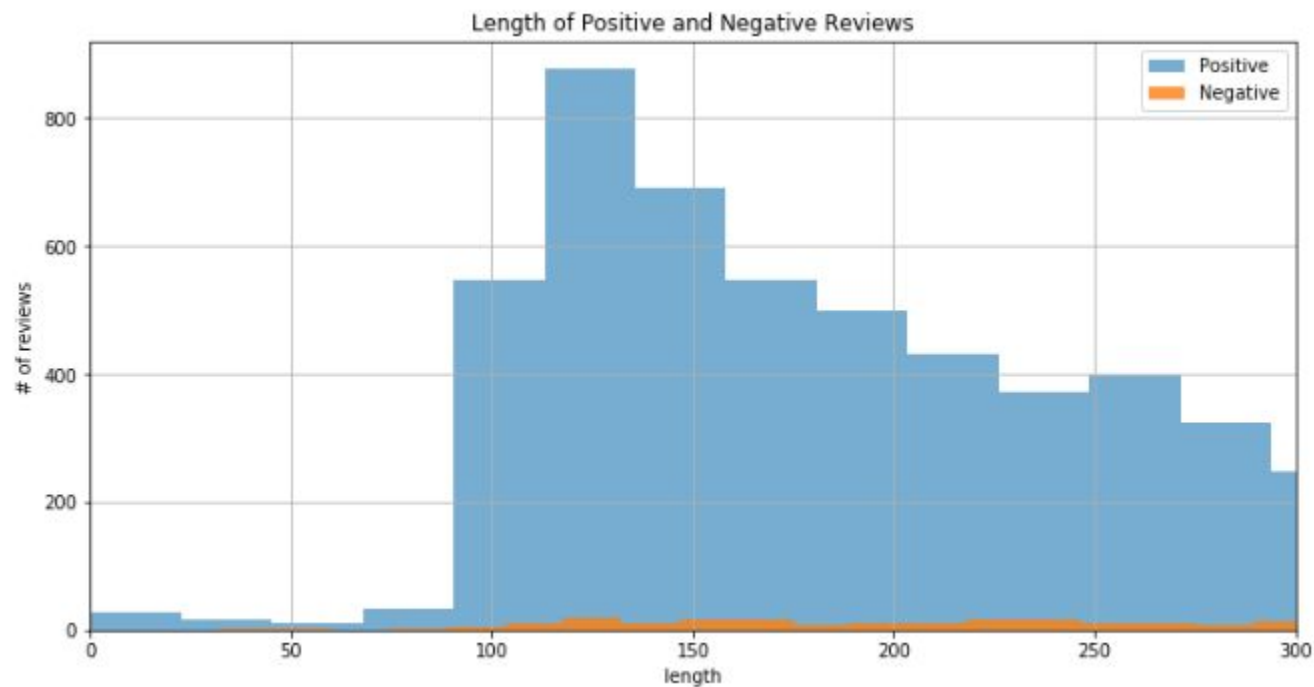
```
Average length of the Postive reviews: 473.7883974216493
Average length of the Negative reviews: 577.7763440860215
```

As you can see the average length of a negative review is longer than a positive review. The average length of the negative review comes in at **578 words** per review, while the positive review comes in at **474 words**.

Lets plot the histogram

```
1 # Lets Look at the distribution of the lengths of the positive versus negative reviews
2
3 import matplotlib
4 from matplotlib import pyplot as plt
5 %matplotlib inline
6
7 matplotlib.rcParams['figure.figsize'] = (12.0, 6.0)
8 bins = 500
9 plt.hist(positive_review['length'], alpha = 0.6, bins=bins, label='Positive')
10 plt.hist(negative_review['length'], alpha = 0.8, bins=bins, label='Negative')
11 plt.xlabel('length')
12 plt.ylabel('# of reviews')
13 plt.title('Length of Positive and Negative Reviews')
14 plt.legend(loc='upper right')
15 plt.xlim(0,300)
16 plt.grid()
17 plt.show()
```


Visualization of the review lengths



Topic modeling

Topic Modeling is a process to automatically identify topics present in a text object and to derive hidden patterns exhibited by a text corpus. Topic Models are very useful for multiple purposes, including:

- Document clustering

- Organizing large blocks of textual data

- Information retrieval from unstructured text

- Feature selection

The goal here is to extract a certain number of groups of important words from the reviews. These groups of words are basically the topics which would help in ascertaining what the consumers are actually talking about in the reviews.

Let's do some data preprocessing

We will remove the punctuations, stopwords and normalize the reviews as much as possible. After every preprocessing step, it is a good practice to check the most frequent words in the data. Therefore, let's define a function that would plot a bar graph of n most frequent words in the data.

```
1 import nltk
2 from nltk import FreqDist
3
4 import gensim
5 from gensim import corpora
6
7 import pyLDAvis
8 import pyLDAvis.gensim


```

```
1 # Lets create a function to plot most frequent terms
2 def freq_words(x, terms = 30):
3     all_words = ' '.join([text for text in x])
4     all_words = all_words.split()
5
6     fdist = FreqDist(all_words)
7     words_df = pd.DataFrame({'word':list(fdist.keys()), 'count':list(fdist.values())})
8
9     # selecting top 20 most frequent words
10    d = words_df.nlargest(columns="count", n = terms)
11    plt.figure(figsize=(20,5))
12    ax = sns.barplot(data=d, x= "word", y = "count")
13    ax.set(ylabel = 'Count')
14    plt.show()

```

The most frequent terms before processing the data

```
1 freq_words(df['reviewText'])
```




As you can see, the most common words are 'the', 'a', 'and', so on. These words are not so important for our task and they do not tell any story. We have to get rid of these kinds of words. Before that let's remove the punctuations and numbers from our text data.

Let's remove the Stopwords

```
1  # Let's try to remove the stopwords and short words (<2 letters) from the reviews.
2
3  from nltk.corpus import stopwords
4  stop_words = stopwords.words('english')
5
6  # function to remove stopwords
7  def remove_stopwords(rev):
8      rev_new = " ".join([i for i in rev if i not in stop_words])
9      return rev_new
10
11 # remove short words (length < 3)
12 df['reviewText'] = df['reviewText'].apply(lambda x: ' '.join([w for w in x.split() if len(w)>2]))
13
14 # remove stopwords from the text
15 reviews = [remove_stopwords(r.split()) for r in df['reviewText']]
16
17 # make entire text lowercase
18 reviews = [r.lower() for r in reviews]
```

After removing the Stowords lets see what the most frequent words now look like.



We can see some improvement here. Terms like 'guitar', 'the', 'one', 'great' have come up which are quite relevant for the Musical equipment category. However, we still have neutral terms like 'the', 'use', 'get', 'also' which are not that relevant.

Let's tokenize the reviews and then lemmatize the tokenized reviews

```
1 tokenized_reviews = pd.Series(reviews).apply(lambda x: x.split())  
2 print(tokenized_reviews[1])
```

```
['the', 'product', 'exactly', 'quite', 'affordable', 'realized', 'double', 'screened', 'arrived', 'even', 'better', 'expected',  
'added', 'bonus', 'one', 'screens', 'carries', 'small', 'hint', 'smell', 'old', 'grape', 'candy', 'used', 'buy', 'reminiscent',  
'sake', 'cannot', 'stop', 'putting', 'pop', 'filter', 'next', 'nose', 'smelling', 'recording', 'dif', 'needed', 'pop', 'filter',  
'work', 'well', 'expensive', 'ones', 'may', 'even', 'come', 'pleasing', 'aroma', 'like', 'mine', 'buy', 'product']
```


```
1 reviews_2 = lemmatization(tokenized_reviews)  
2 print(reviews_2[1]) # print lemmatized review
```

```
['product', 'affordable', 'double', 'screened', 'bonus', 'screen', 'small', 'hint', 'smell', 'old', 'grape', 'candy', 'reminisc  
ent', 'sake', 'pop', 'filter', 'next', 'nose', 'recording', 'dif', 'pop', 'filter', 'work', 'expensive', 'one', 'aroma', 'produ  
ct']
```

We have not just lemmatized the words but also filtered only nouns and adjectives.

Let's de-tokenize the lemmatized reviews and plot the most common words.

After lemmatizing the reviews let take a look at the most frequent word.



It seems that now most frequent terms in our data are relevant. We can now go ahead and start building our topic model.

Using a WordCloud to also visualize the most frequent terms

```
1 # to create our wordclouds, I will import the python module "wordcloud"
2
3 from wordcloud import WordCloud
4
5 # The wordcloud of Reviews in the dataset
6 plt.figure(figsize=(10,7))
7 wordcloud = WordCloud(background_color="black", max_words=10000,
8                       max_font_size= 100)
9 wordcloud.generate(" ".join(df.reviews))
10 plt.title("Keywords in Review", fontsize=20)
11 plt.imshow(wordcloud.recolor(random_state=17), alpha=0.98, interpolation='bilinear')
12 plt.axis('off')
```



LDA Model

Let's start by creating the term dictionary of our corpus, where every unique term is assigned an index

```
1 dictionary = corpora.Dictionary(reviews_2)
2
3 # convert the list of reviews (reviews_2) into a Document Term Matrix using the dictionary prepared above.
4 doc_term_matrix = [dictionary.doc2bow(rev) for rev in reviews_2]
```

```
1 # Creating the object for LDA model using gensim library
2 LDA = gensim.models.ldamodel.LdaModel
3
4 # Build LDA model with 7 topics
5 lda_model = LDA(corpus=doc_term_matrix, id2word=dictionary, num_topics=7, random_state=100)
6
```

Topics that our LDA model has learned

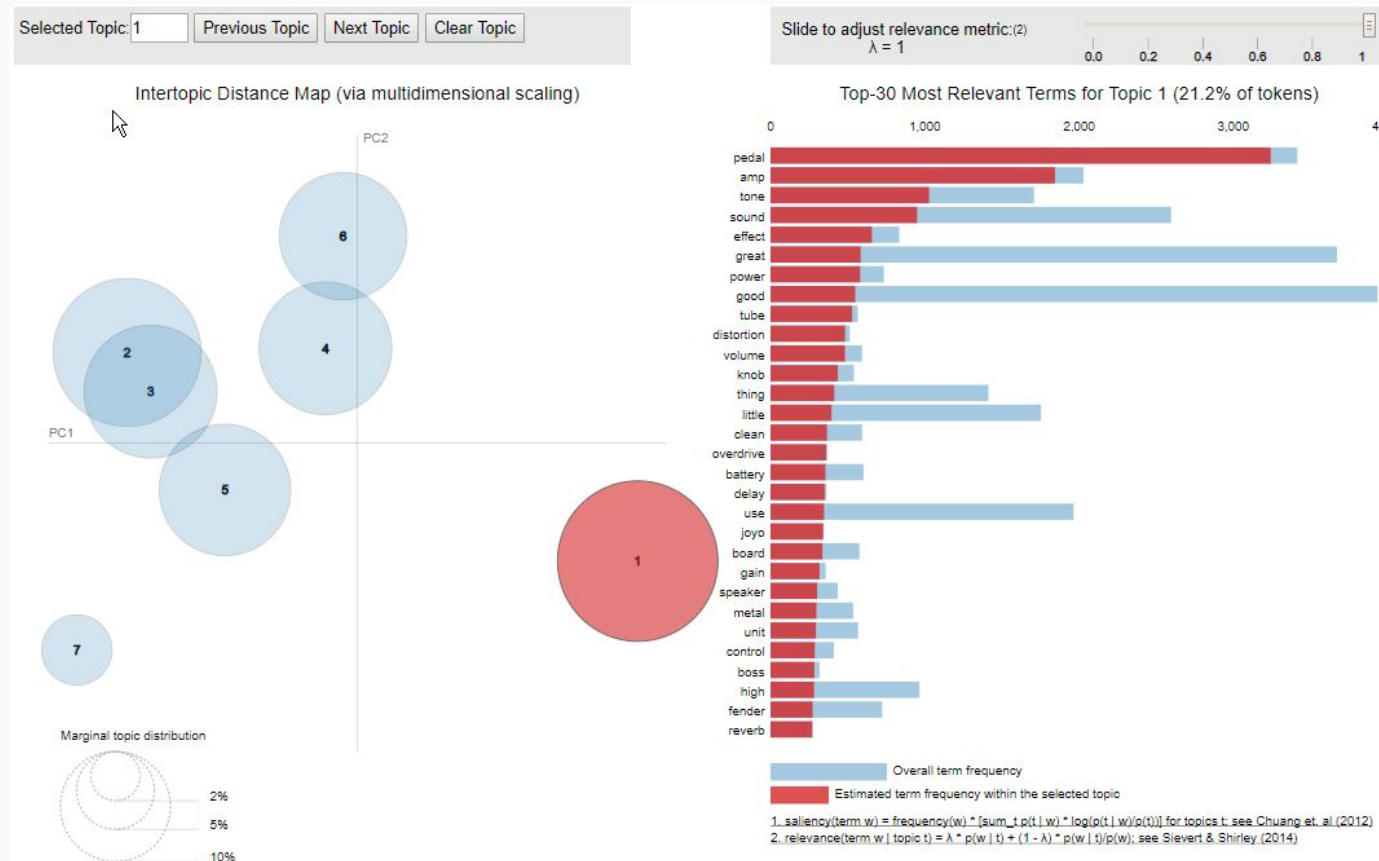
```
1 # Let's print out the topics that our LDA model has learned.  
2  
3 lda_model.print_topics()
```

```
[(0,  
  '0.059*pedal" + 0.033*"amp" + 0.019*"tone" + 0.017*"sound" + 0.012*"effect" + 0.011*"great" + 0.011*"power" + 0.010*"good" +  
  0.010*"tube" + 0.009*"distortion"'),  
 (1,  
  '0.027*"pick" + 0.025*"guitar" + 0.021*"great" + 0.019*"good" + 0.013*"case" + 0.010*"sound" + 0.010*"string" + 0.009*"produc  
  t" + 0.008*"price" + 0.008*"use"'),  
 (2,  
  '0.031*"cable" + 0.029*"great" + 0.025*"good" + 0.019*"price" + 0.019*"quality" + 0.011*"sound" + 0.010*"guitar" + 0.009*"nic  
  e" + 0.008*"product" + 0.008*"use"'),  
 (3,  
  '0.042*"guitar" + 0.038*"tuner" + 0.015*"easy" + 0.014*"strap" + 0.013*"tune" + 0.011*"capo" + 0.011*"use" + 0.010*"good" +  
  0.010*"little" + 0.010*"instrument"'),  
 (4,  
  '0.025*"mic" + 0.014*"microphone" + 0.012*"good" + 0.010*"usb" + 0.009*"use" + 0.009*"great" + 0.008*"display" + 0.008*"recor  
  ding" + 0.008*"sound" + 0.008*"guitar"'),  
 (5,  
  '0.024*"time" + 0.018*"guitar" + 0.015*"little" + 0.014*"violin" + 0.012*"pick" + 0.011*"jam" + 0.010*"nice" + 0.009*"thing"  
  + 0.008*"strap" + 0.008*"mandolin"'),  
 (6,  
  '0.070*"string" + 0.032*"guitar" + 0.020*"good" + 0.013*"sound" + 0.012*"great" + 0.011*"time" + 0.009*"light" + 0.008*"tune"  
  + 0.008*"instrument" + 0.008*"tone"')]
```

Based on topics 1 and 6 terms like guitar, strap and tuner indicate the music instrument that is being reviewed is a guitar. Topics 3 and 5 seems to refer to the overall quality with terms like good and great.

Visualization

To visualize our topics in a 2-dimensional space we will use the pyLDAvis library. This visualization is interactive in nature and displays topics along with the most relevant words



Sentiment Analysis

Sentiment analysis is part of the Natural Language Processing (NLP) techniques that consists in extracting emotions related to some raw texts. This is usually used on customer reviews in order to automatically understand if some users are positive or negative and why.

The goal of this section is to show how sentiment analysis can be performed using python.

Let's create 2 columns for the Sentiment analysis

- Column for sentiments
- Column for number of characters

```
1
2 # add sentiment analysis columns
3 from nltk.sentiment.vader import SentimentIntensityAnalyzer
4
5 sid = SentimentIntensityAnalyzer()
6 df['sentiments'] = df['reviewText'].apply(lambda x: sid.polarity_scores(x))
7 df = pd.concat([df.drop(['sentiments'], axis=1), df['sentiments'].apply(pd.Series)], axis=1)
8
```

```
1 # add number of characters column
2 df['nb_chars'] = df['reviewText'].apply(lambda x: len(x))
3
4 # add number of words column
5 df['nb_words'] = df['reviewText'].apply(lambda x: len(x.split(" ")))
```


Lets create a Doc2vec vector column, train the Doc2vec model, transform each document into a vector data and then add tf-idf column.

```
1 from gensim.test.utils import common_texts
2 from gensim.models.doc2vec import Doc2Vec, TaggedDocument
3
4 documents = [TaggedDocument(doc, [i]) for i, doc in enumerate(df['reviews'].apply(lambda x: x.split(" ")))]
5
6 # Lets train a Doc2Vec model with our text data
7 model = Doc2Vec(documents, vector_size=5, window=2, min_count=1, workers=4)
8
9 # transform each document into a vector data
10 doc2vec_df = df['reviews'].apply(lambda x: model.infer_vector(x.split(" "))).apply(pd.Series)
11 doc2vec_df.columns = ['doc2vec_vector_' + str(x) for x in doc2vec_df.columns]
12 df = pd.concat([df, doc2vec_df], axis=1)
13
```

```
1 # add tf-idfs columns
2 from sklearn.feature_extraction.text import TfidfVectorizer
3 tfidf = TfidfVectorizer(min_df = 10)
4 tfidf_result = tfidf.fit_transform(df['reviews']).toarray()
5 tfidf_df = pd.DataFrame(tfidf_result, columns = tfidf.get_feature_names())
6 tfidf_df.columns = ['word_' + str(x) for x in tfidf_df.columns]
7 tfidf_df.index = df.index
8 df = pd.concat([df, tfidf_df], axis=1)
```

Now lets create a dataset
with the highest positive
sentiment reviews (with
more than 5 words)

```
1 # highest positive sentiment reviews (with more than 5 words)
2 df[df['nb_words'] >= 5].sort_values('pos', ascending = False)[['reviews', 'pos']].head(10)
3
```

	reviews	pos
5945	nice good fine great fantastic love nice good ...	0.956
7509	nice awesome cool ido love	0.755
1481	great great great great great great great grea...	0.746
7605	great strap nice comfort nice pricefit nice fe...	0.717
666	love thing action great sound good easy play r...	0.716
7934	great microphone good rice	0.709
2519	excellent product strong work excellent condit...	0.703
2854	excellent string beautiful solid tone buzzing	0.697
10256	great thank	0.688
1192	easy pack easy instrument light weight sturdy ...	0.683

As expected, the most positive reviews indeed correspond to some good feedbacks.


Now lets create a dataset
with the highest negative
sentiment reviews (with
more than 5 words)

```
1 # lowest negative sentiment reviews (with more than 5 words)
2 df[df['nb_words'] >= 5].sort_values('neg', ascending = False)[['reviews', 'neg']].head(10)
3
```

	reviews	neg
2424	cable less month disappointed quality	0.466
1781	problem able use abuse	0.433
8802	tank couple week time trouble shoot sure compo...	0.433
8553	dull inept version tubescreamer flat	0.397
7165	gig frustrating guess pay	0.387
3697	good nothing special couple review bad smell b...	0.365
9101	cool large inconvenient blue yeti	0.364
5441	good ideal big problem tube amp twin reverb fe...	0.362
9697	dull uninspiring pedal don waste time boss dig...	0.362
3664	bad brother guitar year	0.355

Let's plot the sentiment distribution for positive and negative reviews

```
1 import seaborn as sns
2
3 for x in [0, 1]:
4     subset = df[df['rating'] == x]
5
6     # Draw the density plot
7     if x == 1:
8         label = 'Positive reviews'
9     else:
10        label = 'Negative reviews'
11    sns.distplot(subset['compound'], hist = False, label = label)
```



Predictive Modeling

The dataset is unbalanced and can therefore create false accuracy in the predictive models. To overcome this we can Oversample the data or Undersample the data.

Despite the advantage of balancing classes, these techniques also have their weaknesses. The simplest implementation of over-sampling is to duplicate random records from the minority class, which can cause overfitting. Under-sampling can cause loss of information because we have so few reviews with negative rating.

In this case, we would balance the data with both methods to determine which is most accurate.

Oversampling

```
1 # Determine the max size of the larger rating group
2 max_size = df['rating'].value_counts().max()
```

```
1 max_size
2 print('Max size of the rating:', max_size)
```

Max size of the rating: 8998

```
1 # Lets do some oversampling to take care of the clas imbalance and create a subset.
2 lst = [df]
3 for class_index, group in df.groupby('rating'):
4     lst.append(group.sample(max_size-len(group), replace=True))
5 df_subset = pd.concat(lst)
6 df_subset.head()
```

	asin	helpful	overall	reviewText	reviewTime	reviewerID	reviewerName	summary	unixReviewTime	rating	...
0	1384719342	[0, 0]	5	Not much write about here but does exactly wha...	02/28, 2014	A2IBPI20UZIR0U	cassandra tu "Yeah, well, that's just like, u...	good	1393545600	1	...

The

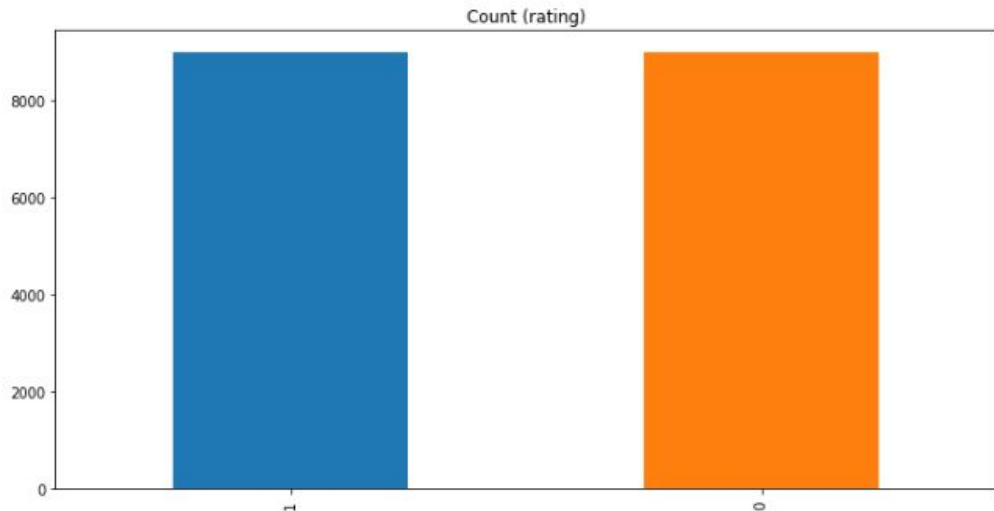
Let's plot the number of positive and negative reviews based on ratings after the Oversampling

```
1 # Number of positive and neative reviews after the oversampling
2 df_subset['rating'].value_counts()
```

```
1    8998
0    8998
Name: rating, dtype: int64
```

```
1 # Lets plot the ratings after the oversampling
2
3 rating_count = df_subset.rating.value_counts()
4
5 print('Negative Rating:', rating_count[0])
6 print('Positive Rating:', rating_count[1])
7 print('Proportion:', round(rating_count[0] / rating_count[1], 2), ': 1')
8
9 rating_count.plot(kind='bar', title='Count (rating)');
```

```
Negative Rating: 8998
Positive Rating: 8998
Proportion: 1.0 : 1
```



Modeling with Multinomial Naive Bayes

In this model we will create a feature using **CountVectorizer**

```
1 # generate document term matrix by using scikit-Learn's CountVectorizer
2
3 from sklearn.feature_extraction.text import CountVectorizer
4 from nltk.tokenize import RegexpTokenizer
5
6 #tokenizer to remove unwanted elements from our data like symbols and numbers
7 token = RegexpTokenizer(r'[a-zA-Z0-9]+')
8 vect = CountVectorizer(lowercase=True, stop_words='english', ngram_range = (1,1), tokenizer = token.tokenize)
9 text_counts= vect.fit_transform(df_subset['reviewText'])
```

```
1 from sklearn.model_selection import train_test_split
2 # Split data into training and test sets
3
4 X_train, X_test, y_train, y_test = train_test_split(text_counts,
5                                                    df_subset['rating'], test_size=0.3, random_state=1)
```

```
1 # fit your model on a train set using fit() and perform prediction on the test set using predict().
2
3 from sklearn.naive_bayes import MultinomialNB
4
5 #Import scikit-learn metrics module for accuracy calculation
6 from sklearn import metrics
7
8 # Model Generation Using Multinomial Naive Bayes
9 clf = MultinomialNB().fit(X_train, y_train)
10 predictions= clf.predict(X_test)
11 print("MultinomialNB Accuracy:", metrics.accuracy_score(y_test, predictions))
12 print('ROC_AUC: ', roc_auc_score(y_test, predictions))
```

```
MultinomialNB Accuracy: 0.9525838118170031
ROC_AUC: 0.952626089629581
```

Modeling with Multinomial Naive Bayes

In this model we will create a feature using **TF-IDF**

```
: 1 from sklearn.feature_extraction.text import TfidfVectorizer
2   tf=TfidfVectorizer()
3   text_tf= tf.fit_transform(df_subset['reviewText'])

: 1 # Split data into training and test sets
2
3   X_train, X_test, y_train, y_test = train_test_split(text_tf,
4                                                         df_subset['rating'], test_size=0.3, random_state=1)

: 1 # Model Building and Evaluation (TF-IDF)
2   clf = MultinomialNB().fit(X_train, y_train)
3   predictions= clf.predict(X_test)
4   print("MultinomialNB Accuracy:",metrics.accuracy_score(y_test, predictions))
5   print('ROC_AUC: ', roc_auc_score(y_test, predictions))
```

```
MultinomialNB Accuracy: 0.9511020559362845
ROC_AUC: 0.9513608704960148
```

Modeling with Logistic Regression

In this model we will create a feature using **CountVectorizer**

```
1 from sklearn.model_selection import train_test_split
2
3 X_train, X_test, y_train, y_test = train_test_split(df_subset['reviewText'],
4                                                    df_subset['rating'], test_size=0.3, random_state=1)
5
```

```
1 vect = CountVectorizer(min_df=5, ngram_range=(1,2)).fit(X_train)
2
3 X_train_vectorized = vect.transform(X_train)
4
5 len(vect.get_feature_names())
```

38792

```
1 from sklearn.linear_model import LogisticRegression
2 from sklearn.metrics import roc_auc_score
3
4 model = LogisticRegression()
5 model.fit(X_train_vectorized, y_train)
6
7 predictions = model.predict(vect.transform(X_test))
8 print("Accuracy:", metrics.accuracy_score(y_test, predictions))
9 print('ROC_AUC: ', roc_auc_score(y_test, predictions))
```

```
Accuracy: 0.9957399518429338
ROC_AUC: 0.9957720588235295
```

As you can see, Logistic regression produces the most accurate model. We are going to stick with the logistic regression model.

Evaluating the logistic regression model with Sensitivity/ Recall and F1 Score.

```
1 # Lets get the F1 score and the recall score
2 from sklearn.metrics import f1_score, recall_score
3 f1_score = round(f1_score(y_test, predictions), 2)
4 recall_score = round(recall_score(y_test, predictions), 2)
5 print("Sensitivity/Recall for Logistic Regression Model 1 : {recall_score}".format(recall_score = recall_score))
6 print("F1 Score for Logistic Regression Model 1 : {f1_score}".format(f1_score = f1_score))
```

```
Sensitivity/Recall for Logistic Regression Model 1 : 0.99
F1 Score for Logistic Regression Model 1 : 1.0
```

Confusion matrix

An interesting way to evaluate the results is by means of a confusion matrix, which shows the correct and incorrect predictions for each class. In the first row, the first column indicates how many classes 0 were predicted correctly, and the second column, how many classes 0 were predicted as 1. In the second row, we note that all class 1 entries were erroneously predicted as class 0.

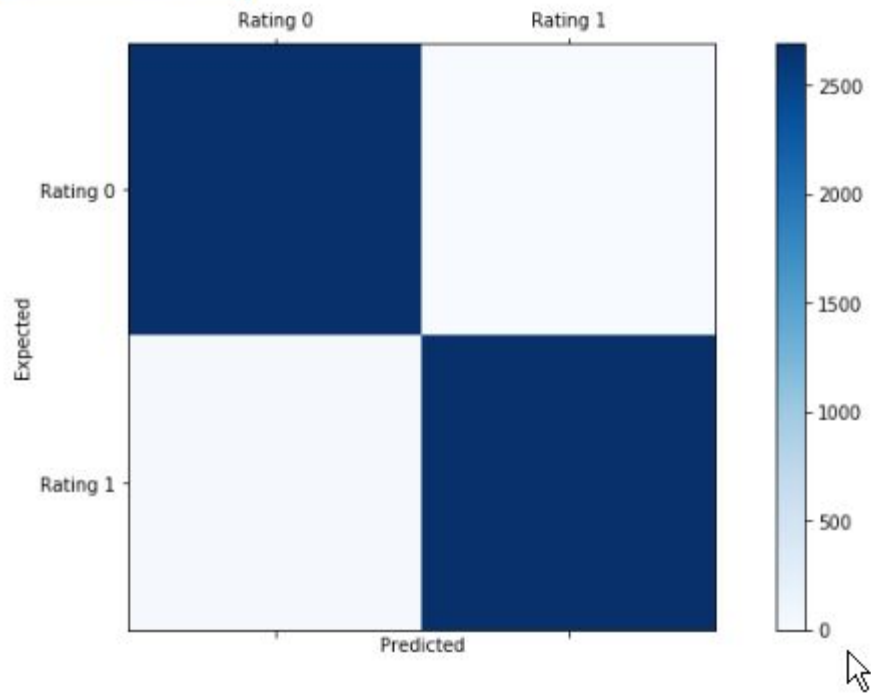
```
1 from sklearn.metrics import confusion_matrix
2 from matplotlib import pyplot as plt
3
4 conf_mat = confusion_matrix(y_true=y_test, y_pred=predictions)
5 print('Confusion matrix:\n', conf_mat)
6
7 labels = ['Rating 0', 'Rating 1']
8 fig = plt.figure()
9 ax = fig.add_subplot(111)
10 cax = ax.matshow(conf_mat, cmap=plt.cm.Blues)
11 fig.colorbar(cax)
12 ax.set_xticklabels([''] + labels)
13 ax.set_yticklabels([''] + labels)
14 plt.xlabel('Predicted')
15 plt.ylabel('Expected')
16 plt.show()
```

Visualization of the Confusion Matrix

The higher the diagonal values of the confusion matrix the better, indicating many correct predictions

Confusion matrix:

```
[[2679  0]  
 [ 23 2697]]
```



Undersampling for comparison

```
1 # Find Number of samples which are negative reviews
2 p_rating = len(df[df['rating'] == 0])
3 # Get indices of positive reviews
4 p_rating_indices = df[df.rating == 1].index
5 # Random sample positive reviews indices
6 random_indices = np.random.choice(p_rating_indices, p_rating, replace=False)
7 # Find the indices of negative reviews samples
8 n_rating_indices = df[df.rating == 0].index
9 # Concat negative review indices with sample positive review ones
10 under_sample_indices = np.concatenate([n_rating_indices, random_indices])
11 # Get Balance Dataframe
12 under_sample = df.loc[under_sample_indices]
13 under_sample.head()
```

	asin	helpful	overall	reviewText	reviewTime	reviewerID	reviewerName	summary	unixReviewTime	rating	length
15	B00005ML71	[0, 0]	2	bought this use with keyboard was really awar...	08 17, 2013	A2PD27UKAD3Q00	Wilhelmina Zeitgeist "coolartsybaby"	Definitely Not For The Seasoned Piano Player	1376697600	0	623
50	B000068NW5	[2, 2]	2	didn't expect this cable thin easily the thickne...	07 6, 2011	A12ABV9NU02O29	C. Longo	Cannot recommend	1309910400	0	387
52	B000068NW5	[0, 0]	1	hums crackles and think having problems with e...	02 9, 2014	A1L7M2JXN4EZCR	David G	I have bought many cables and this one is the ...	1391904000	0	245

Let's plot the
number of positive
and negative
reviews based on
ratings after the
Undersampling

```
1 rating_count_under = under_sample.rating.value_counts()  
2  
3 print('Negative Rating :', rating_count_under[0])  
4 print('Positive Rating :', rating_count_under[1])  
5 print('Proportion:', round(rating_count_under[0] / rating_count_under[1], 2), ': 1')  
6  
7 rating_count_under.plot(kind='bar', title='Count (rating)');
```

Negative Rating : 465
Positive Rating : 465
Proportion: 1.0 : 1



Modeling with Logistic Regression

In this model we will create a feature using **CountVectorizer**

```
1 from sklearn.model_selection import train_test_split
2
3 X_train_under, X_test_under, y_train_under, y_test_under = train_test_split(under_sample['reviewText'],
4                                     under_sample['rating'], test_size=0.3, random_state=1)
5
```

```
1 vect = CountVectorizer(min_df=5, ngram_range=(1,2)).fit(X_train_under)
2
3 X_train_vectorized_under= vect.transform(X_train_under)
4
5 len(vect.get_feature_names())
```

2571

```
1 from sklearn.linear_model import LogisticRegression
2 from sklearn.metrics import roc_auc_score
3 from sklearn import metrics
4
5 model_under = LogisticRegression()
6 model_under.fit(X_train_vectorized_under, y_train_under)
7
8 predictions_under = model_under.predict(vect.transform(X_test_under))
9 print("Accuracy:", metrics.accuracy_score(y_test_under, predictions_under))
10 print('ROC_AUC: ', roc_auc_score(y_test_under, predictions_under))
```

```
Accuracy: 0.7455197132616488
ROC_AUC: 0.7452290076335878
```

Evaluating the logistic regression model with Sensitivity/ Recall and F1 Score.

```
1 from sklearn.metrics import f1_score, recall_score
2 f1_score_under = round(f1_score(y_test_under, predictions_under), 2)
3 recall_score_under = round(recall_score(y_test_under, predictions_under), 2)
4 print("Sensitivity/Recall for Logistic Regression Model 1 : {recall_score_under}".format(recall_score_under = recall_score_under))
5 print("F1 Score for Logistic Regression Model 1 : {f1_score_under}".format(f1_score_under = f1_score_under))
```

```
Sensitivity/Recall for Logistic Regression Model 1 : 0.75
F1 Score for Logistic Regression Model 1 : 0.76
```

Confusion matrix

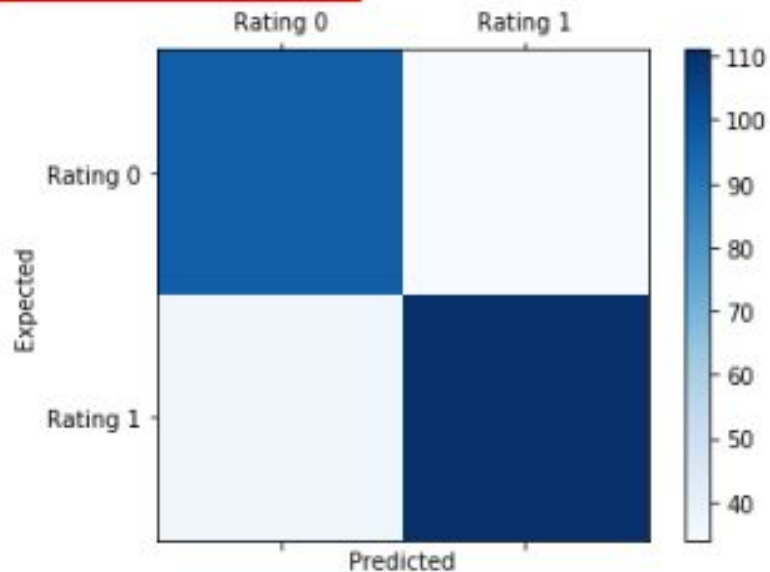
```
1 from sklearn.metrics import confusion_matrix
2 from matplotlib import pyplot as plt
3
4 conf_mat = confusion_matrix(y_true=y_test_under, y_pred=predictions_under)
5 print('Confusion matrix:\n', conf_mat)
6
7 labels = ['Rating 0', 'Rating 1']
8 fig = plt.figure()
9 ax = fig.add_subplot(111)
10 cax = ax.matshow(conf_mat, cmap=plt.cm.Blues)
11 fig.colorbar(cax)
12 ax.set_xticklabels([''] + labels)
13 ax.set_yticklabels([''] + labels)
14 plt.xlabel('Predicted')
15 plt.ylabel('Expected')
16 plt.show()
```


Visualization of the Confusion Matrix

The higher the diagonal values of the confusion matrix the better, indicating many correct predictions

Confusion matrix:

```
[[ 97  34]  
 [ 37 111]]
```



Conclusion

Online product reviews are a great source of information. From the business/sellers' point of view. Online reviews can be used to gauge the consumers' feedback on the products or services they are selling. However, since these online reviews are quite often overwhelming in terms of numbers and information, an intelligent system, capable of finding key insights (topics) from these reviews, will be of great help for both the consumers and the sellers.

We explored the use of data visualization techniques common in data science, such as histograms and pyPlots, to gain a better understanding of the underlying distribution of data in our data set.

As demonstrated above, the dataset had reviews of a musical equipment which included a guitar, microphone and possibly an amp. The sentiment of the the reviews were mostly positive, and the positive reviews were longer in text than negative reviews.

When modeling a dataset with highly unbalanced classes such as this, the classifier always "predicts" the most common class without performing any analysis of the features, which will have a high accuracy rate, which is obviously inaccurate. Oversampling and undersampling was done to determine which would produce a more reliable accuracy rate. As one would predict, oversampling produced the better scores using metrics such as precision, recall, and F1-scores, as well as confusion matrix.

“Information is the oil of the 21st century, and analytics is the combustion engine”

- Peter Sondergaard, Senior Vice President at Gartner

Thanks!

Femi Sulyman

<https://github.com/femi18/Projects>

Femisu@gmail.com

